

# Optimising the LINPACK 1000 using Parallelization

Sam Serrels  
40082367@napier.ac.uk  
Edinburgh Napier University  
Concurrent and Parallel Systems (SET10108)

## 1 Introduction

**The LINPACK benchmark** The Linpack algorithm is a popular benchmark in the high performance computing field as it uses as a floating point performance measure for ranking supercomputers. The Benchmark solves a large system of linear equations using LU decomposition and is typical of many matrix-based scientific computations. Linpack started as a Fortran maths processing application, the code for solving linear equations was extracted from the original program and turned into a benchmark.

The Linpack1000 algorithm first generates a random 1000x1000 element matrix, A, and 1000 element vector, b. The elements in A and b are all double precision floating point numbers. Processing then takes place to find the solution, a 1000 element vector, x, such that  $Ax = b$ .

This report documents the process of analysing a specific sequential Linpack implementation and then converting it to a parallel task.

**Related Work** As Linpack is used to benchmark large scale multi-processor supercomputers, there are many parallel versions available. The main difference between implementations is the Gaussian elimination stage, for which there are many algorithms available, some of which can split the task up into separate easily parallelizable chunks of logic. For the scope of this project, changes to the algorithm were avoided wherever possible to keep a fair comparison to the original sequential code.

**High-Performance Linpack** The most commonly used implementation of Linpack is the HPL implementation, written by the Innovative Computing Laboratory at the University of Tennessee. HPL is an open-source project that aims to provide a toolbox for configuring, optimising, and running the benchmark over a network. It contains many versions of the algorithm with plenty of configurable options and for tuning performance to a specific system.

The HPL project was used as a rough guide to how the reference algorithm could be modified for this project, however most of the optimisation were beyond the scope of this project.

**Project Scope** A quick optimisation would be to drop the precision of the algorithm from double, to single precision floating point values. Another method would be to swap the original Gaussian elimination algorithm for a different mathematical approach that would lend itself to parallel processing better. This project aimed to see how much the original Linpack code can be optimised with parallelization, without changing the core logic of the algorithm or data output so these routes for optimisation were ruled out.

**OpenMP** The technology for processing the application in parallel was chosen to be OpenMp, an API that abstracts the creation of threads from the user and therefore allows for easier development and better cross platform portability, assuming that the chosen platform has a compiler that supports OpenMP. This was chosen over creating threads manually, mainly for ease of development reasons, but also because even in a situation that OpenMp is slower than Manual threads, there should still be a noticeable performance increase over the baseline results.

**SIMD** For an extra level of performance, SIMD instructions were used to gain performance in the most frequently executed parts of the

program.

## 2 Linpack Gaussian elimination




Figure 1: Overview of GE progression -




Figure 2: Total program execution -

**Initial analysis** The Gaussian elimination(GE) stage, work it's way through the array, starting in the "top left corner". It examines the entire first column(C) ("Pivot and scale", figure ??) and finds the largest value (T). The row that contains the largest value(T) becomes the pivot, it is swapped with the topmost row. Then, each column is processed, by multiplying it by  $1/T$  and then adding the value of the C column.

Once this is complete, the process restarts but in a subsection of the

Matrix A that is one row and column smaller than the previous iteration. This continues until the "Bottom right" corner is reached. This process transforms Matrix A into an upper triangular matrix that is in row echelon form.

On analysis of execution of the program, it was clear that the vast majority of the execution time was based in the Gaussian Elimination stage. On further examination, the "Daxpy" function, which is called many times during the GE stage takes up nearly 90% of the total execution time. Daxpy is the function that does the scaling and addition of each column, and is called 424166 times during the full execution of the program.

**Daxpy** While the Daxpy function is called at a high frequency it contains very few lines of code. Its function is to compute  $Y = S * X + Y$ , X and Y are elements of two arrays, and S is a scaler value. In the program, this is used in a loop to process each column in the A array.

This was the first part of the code to be examined for possible speed-up, as each iteration of the loop doesn't depend on any other iteration. Initially Parallelizing the loop with OpenMp was attempted, but as the loop will only ever execute a maximum of 1000 times, the overhead time of creating and running threads was always greater than the time of running the loop without threads.

**Listing 1: daxpy Code**

```
1 void (int n, double scaler, double *dx, double *dy, int offset) {
2     double *const y = &dy[offset];
3     double *const x = &dx[offset];
4     for (int i = 0; i < n; ++i) {
5         y[i] += scaler * x[i];
6     }
7 }
8
```

**Simd Daxpy** As no increase of performance could be gained by using more threads, Simd instructions were investigated. Daxpy processes N amount of numbers from the two input arrays, using a loop. The numbers to be processed are sequentially laid out in the input arrays, so the logic to convert the loop from processing one item at a time, to multiple, was a simple task. Additional lines of code were unavoidable, e.g if the size of N was not divisible by the amount of items that the loop processes in one pass, then the remainder would have to be processed at the end.

Simd instructions requires aligned memory, therefore the numbers from the input arrays had to be loaded into special aligned containers. This could be avoided if the whole program was converted to use aligned data arrays.

**Listing 2: Simd daxpy Code**

```
1 void (int n, double scaler, double *dx, double *dy, int offset) {
2     if ((n <= 0) || (scaler == 0)) {
3         return;
4     }
5     double *const y = &dy[offset];
6     double *const x = &dx[offset];
7
8     const __m256d scalars = _mm256_set1_pd(scaler);
9     const int remainder = n % 4;
10    const int nm1 = n - 3;
11
12    for (int i = 0; i < nm1; i += 4) {
13        // load X
14        const __m256d xs = _mm256_loadu_pd(&x[i]);
15        // load y
16        __m256d ys = _mm256_loadu_pd(&y[i]);
17        // multiply X by scalars, add to Y
18        ys = _mm256_add_pd(ys, _mm256_mul_pd(xs, scalars));
19        // load back into y
```

```
20        __mm256_storeu_pd(&y[i], ys);
21    }
22
23    for (int i = n - remainder; i < n; ++i) {
24        y[i] += scaler * x[i];
25    }
26 }
```

**Simd Daxpy Results** Both 128bit and 256bit (Largest supported by available hardware) Simd versions of Daxpy were implemented and tested. The 128bit version provided no performance gains, due to the overhead of converting input data to aligned data. 256 bit instructions however provided a significant increase in performance, but only when used with an N parameter greater than around 100. As seen in Figure ??.

#### The Pivot loop

#### The Column loop

#### Memory alignment

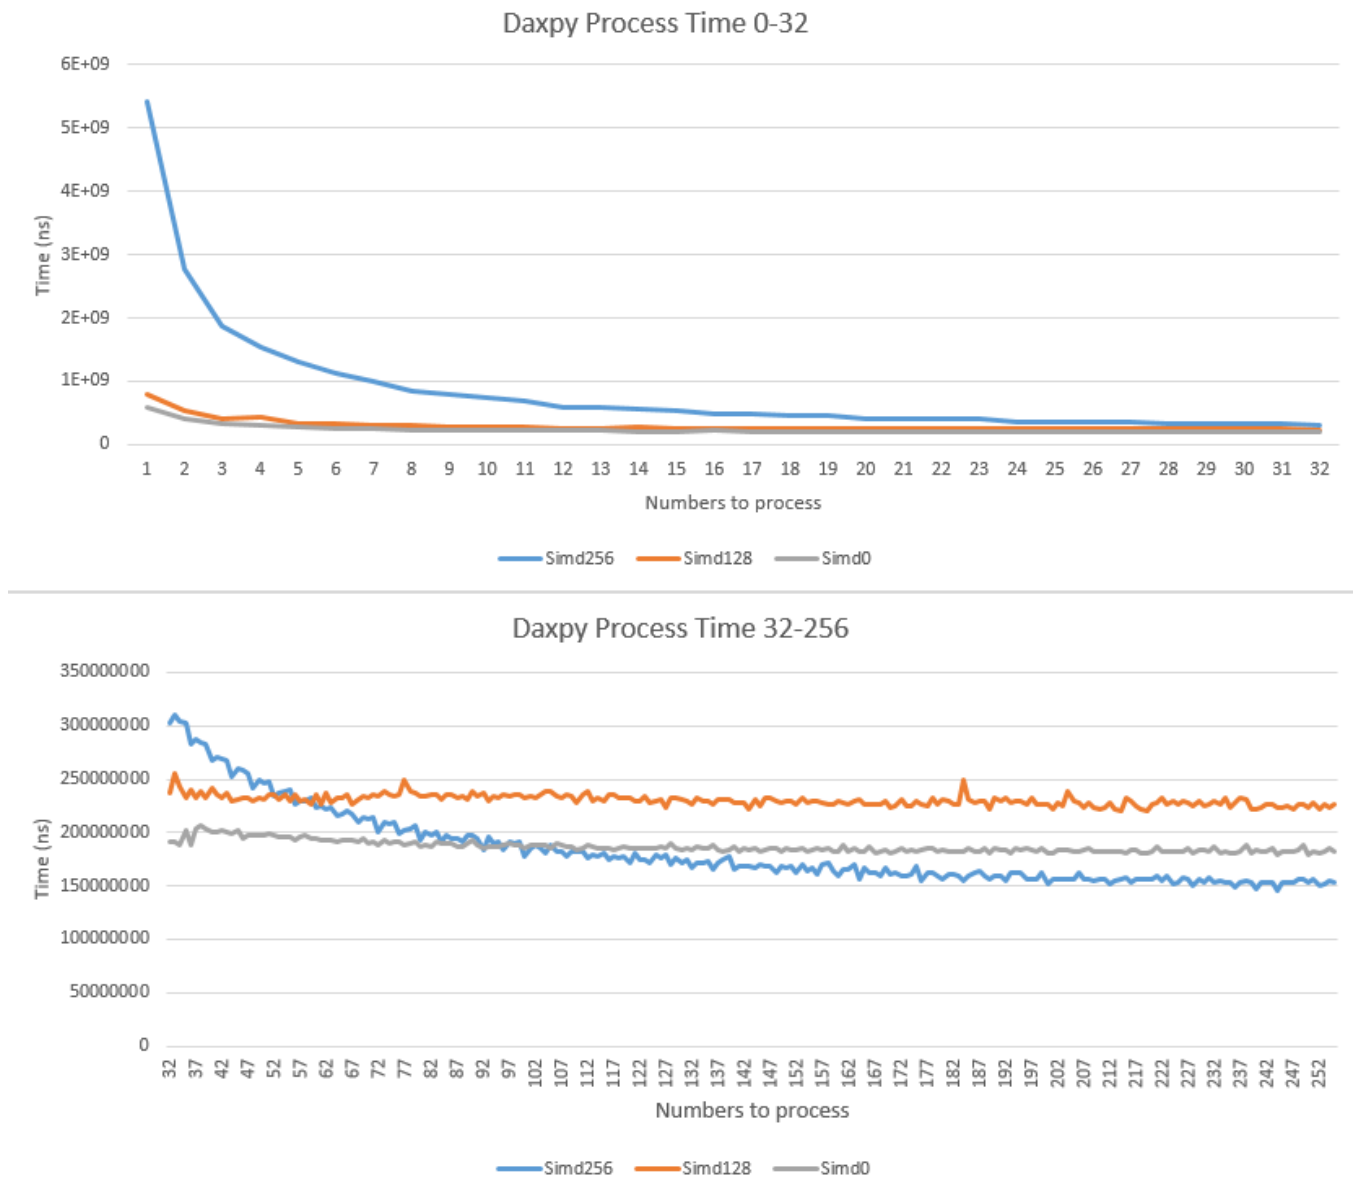
### 3 Optimisation Method

#### Code Simplification

**Parallelised Daxpy** no route was over the capacity of the truck.

#### SIMD Daxpy

#### MDaxpy



**Figure 3: Daxpy Simd Comparisons** - Time taken to process 10'000 numbers, 10'000 times Each subsequent call to Daxpy increases the amount of number to calculate at once.

Name	Allocate Memory (ms)	Create Input Numbers (ms)	gaussian eliminate (ms)	Solve (ms)	Validate (ms)	Total Time (ms)	Total Speedup	Speedup (With Simd)*
Threads: 1 No Simd	0.48	5.22	157.27	0.69	5.27	168.93	0%	0%
Threads: 1 Simd128	0.45	5.00	152.68	0.67	5.06	163.85	3%	0%
Threads: 1 Simd256	0.44	5.06	142.04	0.67	5.14	153.34	9%	0%
Threads: 2 No Simd	0.42	5.48	78.76	0.68	5.52	90.86	46%	46%
Threads: 2 Simd128	0.46	5.36	85.35	0.66	5.49	97.32	42%	41%
Threads: 2 Simd256	0.45	4.82	72.86	0.63	4.83	83.59	51%	45%
Threads: 3 No Simd	0.45	5.37	63.49	0.68	5.46	75.44	55%	55%
Threads: 3 Simd128	0.46	5.52	60.00	0.66	5.64	72.28	57%	56%
Threads: 3 Simd256	0.46	5.74	52.51	0.58	5.85	65.14	61%	58%
Threads: 4 No Simd	0.44	5.28	42.67	0.62	5.33	54.34	68%	68%
Threads: 4 Simd128	0.43	5.72	41.51	0.54	5.78	53.97	68%	67%
Threads: 4 Simd256	0.48	5.80	39.61	0.55	5.91	52.36	69%	66%
Threads: 5 No Simd	0.46	5.41	50.89	0.67	5.50	62.93	63%	63%
Threads: 5 Simd128	0.45	5.20	49.49	0.67	5.24	61.05	64%	63%
Threads: 5 Simd256	0.46	5.11	46.46	0.60	5.24	57.86	66%	62%
Threads: 6 No Simd	0.50	5.80	44.44	0.75	5.93	57.41	66%	66%
Threads: 6 Simd128	0.48	5.53	43.04	0.69	5.66	55.41	67%	66%
Threads: 6 Simd256	0.49	5.55	41.63	0.65	5.72	54.04	68%	65%
Threads: 7 No Simd	0.53	6.00	40.21	0.75	6.18	53.66	68%	68%
Threads: 7 Simd128	0.52	5.40	40.02	0.73	5.55	52.22	69%	68%
Threads: 7 Simd256	0.55	5.58	35.69	0.68	5.77	48.28	71%	69%
Threads: 8 No Simd	0.51	5.01	43.25	0.80	5.17	54.75	68%	68%
Threads: 8 Simd128	0.53	5.48	43.44	0.78	5.73	55.96	67%	66%
Threads: 8 Simd256	0.51	5.47	38.64	0.71	5.72	51.05	70%	67%

**Table 1:** Results of all tests

*\*Speed-up with Simd, compares times against the Simd equivalent 1 Thread run*




Figure 4: Total Speed-up percentage, for each number of threads -




Figure 5: Total Speed-up (Relative to baseline Simd) percentage, for each number of threads -

## 4 Conclusions

## References


CLARKE, G., AND WRIGHT, J. 1962. Scheduling of vehicles from a central depot to a number of delivery points. *Operations Research* 12, 4, 568–581.

**Computation time** Both Implementations of the Clarke-Wright algorithms produced expected results, with the parallel version producing larger and fewer routes. As for the time taken to calculate, the performance is roughly equal. The discrepancies shown in Figure ?? when the amount of customers increases beyond 800 is possibly due to optimisations carried out by the Java virtual machine. The total operations carried out is roughly the same for each algorithm, however the arrays are accessed and modified at different times, this is a possible cause for the difference in processing time.


**Solution Quality** The Parallel solution produced a large saving of up to a 600% increase against the baseline cost, shown in Figure ?. The Sequential solution produced a constant saving of around 200%. These results are also shown in Figure ?, showing the number of routes.

**Edge Cases** It is possible that a customer can be left out of all routes due to capacity constraints; this is checked for at the end of the calculation. If a customer is left over, it is seen if it would be possible to add it to any existing route and then if it would be more efficient than sending out a new truck. This can be seen in Figure ?, the customer in the top left falls in to this edge case category.


**Conclusion** The implementation written for this report successfully computes optimised and usable data, the processing cost increases in a quadratic relation to the size of data. The specific implementation could be optimised to produce quicker results. One possible optimisation route could be a custom sort method, as profiling reported that 40% of the processing time is taken by the initial sort of the customer pairs. Overall this report produced repeatable and meaningful data, and can be seen as a successful investigation into the Clark-Wright Algorithm.




**Figure 6:** *Single Threaded, 6 runs, simd256 Daxpy - Overall system CPU utilisation*




**Figure 7:** *Single Threaded, 6 runs, simd256 Daxpy - Thread to CPU Core allocation*




**Figure 8:** *Single Threaded, 6 runs, simd256 Daxpy - Thread Status*



**Figure 9: 4 Threads, 6 runs, simd256 Daxpy - Overall system CPU utilisation**




**Figure 10: 4 Threads, 6 runs, simd256 Daxpy - Thread to CPU Core allocation**




**Figure 11: 4 Threads, 6 runs, simd256 Daxpy - Thread Status**






**Figure 12:** 8 Threads, 6 runs, *simd256 Daxpy* - Overall system CPU utilisation



**Figure 13:** 8 Threads, 6 runs, *simd256 Daxpy* - Thread to CPU Core allocation



**Figure 14:** 8 Threads, 6 runs, *simd256 Daxpy* - Thread Status

## 5 Appendix: Code

```
1 #include <iostream>
2 #include <string>
3 #include <assert.h>
4 #include <omp.h>
5 #include "sequentialOMP.h"
6 #include "Timer.h"
7 #include <algorithm>
8 #include <cvmarkersobj.h>
9
10 #define MAX_THREADS 8
11 #define PAR_DAXPY 0
12 #define SIMD_DAXPY256 0
13 #define SIMD_DAXPY128 0
14 #define SIMD_Loop 1
15 using namespace std;
16 using namespace Concurrency::diagnostic;
17
18 namespace seqOMP {
19
20 void (*cdaxpy)(unsigned int, const double, double *, double *, ←
    unsigned int);
21 int (*cgaussian)(double **, int, int *);
22 // Fills A with random doubles, populates b with row sums, returns ←
    largest val
23 double fillArray(double **a, int n, double *b) {
24     double largestValue = 0.0;
25     int init = 1325;
26
27     for (int i = 0; i < n; ++i) {
28         for (int j = 0; j < n; ++j) {
29             init = 3125 * init % 65536; // cheap and nasty random generator
30             a[j][i] = (static_cast<double>(init) - 32768.0) / 16384.0;
31             largestValue = (a[j][i] > largestValue) ? a[j][i] : largestValue;
32         }
33     }
34     // fill b with 0
35     for (int i = 0; i < n; ++i) {
36         b[i] = 0.0;
37     }
38     // add every element of each row of A to each row of B
39     for (int j = 0; j < n; ++j) {
40         for (int i = 0; i < n; ++i) {
41             b[i] += a[j][i];
42         }
43     }
44
45     return largestValue;
46 }
47
48 /* Purpose : Find largest component of double vector dx
49 n : number of elements in input vector
50 dx : double vector with n+1 elements, dx[0] is not used
51 dx_off : offset in reading dx
52 */
53 int indexOfLargestElement(int n, double *dx, int dx_off) {
54     double dmax, dtemp;
55     int itemp = 0;
56
57     if (n < 1) {
58         itemp = -1;
59     } else if (n == 1) {
60         itemp = 0;
61     } else {
62         itemp = 0;
63         dmax = abs(dx[0 + dx_off]);
64         for (int i = 0; i < n; ++i) {
65             dtemp = abs(dx[i + dx_off]);
66             if (dtemp > dmax) {
67                 itemp = i;
68                 dmax = dtemp;
69             }
70         }
71     }
72
73     return itemp;
74 }
75
76 // Scales a vector by a constant
77 void scaleVecByConstant(int n, double da, double *dx, int dx_off, int ←
    incx) {
78     if (n > 0) {
79         if (incx != 1) {
80             int nincx = n * incx;
81             for (int i = 0; i < nincx; i += incx) {
82                 dx[i + dx_off] *= da;
83             }
84         } else {
85             for (int i = 0; i < n; ++i) {
86                 dx[i + dx_off] *= da;
87             }
88         }
89     }
90     /* Constant times a vector plus a vector
91     Purpose : To compute dy = da * dx + dy
92     ---- Input ----
93     n : number of elements in input vector(s)
94     scaler : double scalar multiplier
95     dx : double vector with n+1 elements
96     dy : double vector with n+1 element
97     ---- Output ----
98     dy = da * dx + dy, unchanged if n <= 0
99 */
100 void daxpyS128(unsigned int n, const double scaler, double *dx, ←
    double *dy,
    unsigned int offset) {
101
```

```

102 if ((n <= 0) || (scaler == 0)) {
103     return;
104 }
105 double *const y = &dy[offset];
106 double *const x = &dx[offset];
107 const _mm128d scalars = _mm_set1_pd(scaler);
108 const int remainder = n % 2;
109 const int nm1 = n - 1;
110
111 for (int i = 0; i < nm1; i += 2) {
112     // load X
113     const _mm128d xs = _mm_loadu_pd(&x[i]);
114     // load y
115     _mm128d ys = _mm_loadu_pd(&y[i]);
116     // multiply X by scalars, add to Y
117     ys = _mm_add_pd(ys, _mm_mul_pd(xs, scalars));
118     // load back into y
119     _mm_storeu_pd(&y[i], ys);
120 }
121
122 if (remainder != 0) {
123     y[n - 1] += scaler * x[n - 1];
124 }
125 }
126 void daxpyS256(unsigned int n, const double scaler, double *dx, ←
    double *dy,
    unsigned int offset) {
127     if ((n <= 0) || (scaler == 0)) {
128         return;
129     }
130 }
131 double *const y = &dy[offset];
132 double *const x = &dx[offset];
133
134 const _mm256d scalars = _mm256_set1_pd(scaler);
135 const int remainder = n % 4;
136 const int nm1 = n - 3;
137
138 for (int i = 0; i < nm1; i += 4) {
139     // load X
140     const _mm256d xs = _mm256_loadu_pd(&x[i]);
141     // load y
142     _mm256d ys = _mm256_loadu_pd(&y[i]);
143     // multiply X by scalars, add to Y
144     ys = _mm256_add_pd(ys, _mm256_mul_pd(xs, scalars));
145     // load back into y
146     _mm256_storeu_pd(&y[i], ys);
147 }
148
149 for (int i = n - remainder; i < n; ++i) {
150     y[i] += scaler * x[i];
151 }
152 }
153 void daxpyPar(unsigned int n, const double scaler, double *dx, ←
    double *dy,
    unsigned int offset) {
154     if ((n <= 0) || (scaler == 0)) {
155         return;
156     }
157 }
158 double *const y = &dy[offset];
159 double *const x = &dx[offset];
160 #pragma omp parallel
161 {
162     const int thread_id = omp_get_thread_num();
163     const int thread_count = omp_get_num_threads();
164     const int per_thread = n / thread_count;
165
166     for (int i = (thread_id * per_thread); i < per_thread; ++i) {
167         y[i] += scaler * x[i];
168     }
169 }
170 }
171
172 void daxpy(unsigned int n, const double scaler, double *dx, double *←
    dy,
    unsigned int offset) {
173     double *const y = &dy[offset];
174     double *const x = &dx[offset];
175     for (int i = 0; i < n; ++i) {
176         y[i] += scaler * x[i];
177     }
178 }
179 }
180 }
181
182 // Performs Gaussian elimination with partial pivoting
183 int gaussian_eliminate(double **a, int n, int *ipivot) {
184     // Pointers to columns being worked on
185     double *col_k;
186     int nm1 = n - 1;
187     int info = 0;
188
189     if (nm1 >= 0) {
190         int kp1, l;
191         for (int k = 0; k < nm1; ++k) {
192             // Set pointer for col_k to relevant column in a
193             col_k = &a[k][0];
194             kp1 = k + 1;
195
196             // Find pivot index
197             l = index_of_largest_element(n - k, col_k, k) + k;
198             ipivot[k] = l;
199
200             // Zero pivot means that this column is already triangularized
201             if (col_k[l] != 0) {
202                 double t;
203                 // Check if we need to interchange
204                 if (l != k) {
205                     t = col_k[l];
206                     col_k[l] = col_k[k];
207                     col_k[k] = t;
208                 }
209
210                 // Compute multipliers
211                 t = -1.0 / col_k[k];
212                 // Multiply column by t
213                 scale_vec_by_constant(n - kp1, t, col_k, kp1, 1);
214
215                 // Row elimination with column indexing
216                 for (int j = kp1; j < n; ++j) {
217                     // Set pointer for col_j to relevant column in a
218                     double *col_j = &a[j][0];
219
220                     double t = col_j[l];
221                     if (l != k) {
222                         col_j[l] = col_j[k];
223                         col_j[k] = t;
224                     }
225                     daxpy(n - kp1, t, col_k, col_j, kp1);
226                 }
227             } else
228                 info = k;
229         }
230     }
231 }
232
233 ipivot[n - 1] = n - 1;
234 if (a[n - 1][n - 1] == 0) {
235     info = n - 1;
236 }
237
238 return info;
239 }
240 // Performs Gaussian elimination with partial pivoting
241 int gaussian_eliminatePAR(double **a, int n, int *ipivot) {
242     // Pointers to columns being worked on
243     double *col_k;
244     int nm1 = n - 1;
245     int info = 0;
246
247     if (nm1 >= 0) {
248         int kp1, l;
249         for (int k = 0; k < nm1; ++k) {
250             // Set pointer for col_k to relevant column in a
251             col_k = &a[k][0];
252             kp1 = k + 1;
253
254             // Find pivot index
255             l = index_of_largest_element(n - k, col_k, k) + k;
256             ipivot[k] = l;

```

```

257 // Zero pivot means that this column is already triangularized
258 if (col_k[l] != 0) {
259     double t;
260     // Check if we need to interchange
261     if (l != k) {
262         t = col_k[l];
263         col_k[l] = col_k[k];
264         col_k[k] = t;
265     }
266
267     // Compute multipliers
268     t = -1.0 / col_k[k];
269     scaleVecByConstant(n - kp1, t, col_k, kp1, 1);
270
271 // Row elimination with column indexing
272 #pragma omp parallel for
273 for (int j = kp1; j < n; ++j) {
274     // Set pointer for col_j to relevant column in a
275     double *col_j = &a[j][0];
276
277     double t = col_j[l];
278     if (l != k) {
279         col_j[l] = col_j[k];
280         col_j[k] = t;
281     }
282     cdaxpy(n - kp1, t, col_k, col_j, kp1);
283 }
284 } else
285     info = k;
286 }
287
288 ipivot[n - 1] = n - 1;
289 if (a[n - 1][n - 1] == 0) {
290     info = n - 1;
291 }
292 return info;
293 }
294
295 // gaussian_eliminate
296 void dgesl(double **a, int n, int *ipivot, double *b) {
297     int k, nm1;
298     nm1 = n - 1;
299
300 // Solve a * x = b. First solve l * y = b
301 if (nm1 >= 1) {
302     for (k = 0; k < nm1; ++k) {
303         int l = ipivot[k];
304         double t = b[l];
305
306         if (l != k) {
307             b[l] = b[k];
308             b[k] = t;
309         }
310
311         cdaxpy(n - (k + 1), t, &a[k][0], b, (k + 1));
312     }
313
314 // Now solve u * x = y
315 for (int kb = 0; kb < n; ++kb) {
316     k = n - (kb + 1);
317     b[k] /= a[k][k];
318     double t = -b[k];
319     cdaxpy(k, t, &a[k][0], b, 0);
320 }
321
322 // Multiply matrix m times vector x and add the result to vector y
323 void dmxpy(int n1, double *y, int n2, double *x, double **m) {
324     for (int j = 0; j < n2; ++j) {
325         for (int i = 0; i < n1; ++i) {
326             y[i] += x[j] * m[j][i];
327         }
328     }
329 }
330
331 }
332
333 }
334
335 }
336 }
337 }
338
339 // Runs the benchmark
340 void run(double **a, double *b, int n, int *ipivot) {
341
342 // Validates the result
343 void validate(double **a, double *b, double *x, int n) {
344     // copy b into x
345     for (int i = 0; i < n; ++i) {
346         x[i] = b[i];
347     }
348
349 // reset A and B arrays to original rand values
350 double biggestA = fillArray(a, n, b);
351
352 for (int i = 0; i < n; ++i) {
353     b[i] = -b[i];
354 }
355
356 // multiply a*x, add to b
357 dmxpy(n, b, n, x, a);
358
359 double biggestB = 0.0;
360 double biggestX = 0.0;
361 for (int i = 0; i < n; ++i) {
362     biggestB = (biggestB > abs(b[i])) ? biggestB : abs(b[i]);
363     biggestX = (biggestX > abs(x[i])) ? biggestX : abs(x[i]);
364 }
365
366 double residn =
367     biggestB / (n * biggestA * biggestX * (2.2204460492503131e-16));
368 assert(residn < CHECK_VALUE);
369 }
370
371 int start(const unsigned int runs, const unsigned int threadCount,
372     const bool simd128, const bool simd256) {
373
374     double *a = new double [10000];
375     double *b = new double [10000];
376     double *c = new double [10000];
377
378     int init = 1325;
379     for (int j = 0; j < 10000; ++j) {
380         init = 3125 * init % 65536; // cheap and nasty random generator
381         a[j] = (static_cast<double>(init) - 32768.0) / 16384.0;
382     }
383     for (int j = 0; j < 10000; ++j) {
384         init = 3125 * init % 65536; // cheap and nasty random generator
385         b[j] = (static_cast<double>(init) - 32768.0) / 16384.0;
386     }
387     for (int j = 0; j < 10000; ++j) {
388         init = 3125 * init % 65536; // cheap and nasty random generator
389         c[j] = (static_cast<double>(init) - 32768.0) / 16384.0;
390     }
391
392     for (int q = 1; q < 256; ++q) {
393         auto start = chrono::high_resolution_clock::now();
394         for (int j = 0; j < 10000; ++j) {
395             for (int i = 0; i < 10000; i += q) {
396                 daxpy(q, c[i], a, b, i);
397             }
398         }
399     }
400     cout << q << ",\t" << chrono::duration_cast<chrono::nanoseconds>(chrono::high_resolution_clock::now() - start).count() << endl;
401 }
402
403 return 0;
404
405 ResultFile r;
406 r.name = "Sequential LinPack OMP" + to_string(runs);
407
408 if (simd256) {
409     cdaxpy = &daxpyS256;
410     r.name += "Simd256";
411 }

```

```

412 } else if (simd128) {
413     cdaxpy = &daxpyS128;
414     r.name += "Simd128";
415 } else {
416     cdaxpy = &daxpy;
417 }
418 if (threadCount > 1) {
419     unsigned int t = min((unsigned int)omp_get_max_threads(), ←
        threadCount);
420     cgaussian = &gaussian_eliminatePAR;
421     omp_set_num_threads(t);
422     r.name += "ParLoopT" + to_string(t);
423 } else {
424     cgaussian = &gaussian_eliminate;
425 }
426 cout << r.name << endl;
427 r.headings = {"Allocate Memory", "Create Input Numbers",
428     "gaussian_eliminate", "Solve", "Validate"};
429
430 marker_series series;
431 // span *flagSpan = new span(series, 1, _T("flag span"));
432 // series.write_flag(_T("Here is the flag.));
433 // delete flagSpan;
434
435 Timer time_total;
436 for (size_t i = 0; i < runs; i++) {
437     // series.write_flag(0, (to_string(i).c_str()));
438     series.write_flag(_T("Here is the flag.));
439     span *flagSpan = new span(series, 1, _T("time_allocate"));
440     cout << i << endl;
441     // Allocate data on the heap
442     Timer time_allocate;
443     double **a = new double *[NSIZE];
444     for (size_t i = 0; i < NSIZE; ++i) {
445         // a[i] = new double[SIZE];
446         a[i] = (double *)_aligned_malloc(NSIZE * sizeof(double), sizeof(←
            double));
447     }
448
449     double *b =
450         (double *)_aligned_malloc(NSIZE * sizeof(double), sizeof(←
            double));
451
452     double *x =
453         (double *)_aligned_malloc(NSIZE * sizeof(double), sizeof(←
            double));
454
455     int *ipivot = (int *)_aligned_malloc(NSIZE * sizeof(int), sizeof(int)←
        );
456     // double *b = new double[SIZE];
457     // double *x = new double[SIZE]
458     ;
459     // int *ipivot = new int[SIZE];
460     delete flagSpan;
461     time_allocate.Stop();
462
463     // Main application
464     flagSpan = new span(series, 1, _T("time_genRnd"));
465     Timer time_genRnd;
466
467     auto aa = fillArray(a, NSIZE, b);
468
469     delete flagSpan;
470     time_genRnd.Stop();
471
472     flagSpan = new span(series, 1, _T("time_gauss"));
473     Timer time_gauss;
474
475     cgaussian(a, NSIZE, ipivot);
476
477     delete flagSpan;
478     time_gauss.Stop();
479
480     flagSpan = new span(series, 1, _T("solve"));
481     Timer time_dgesl;
482
483     dgesl(a, NSIZE, ipivot, b);
484
485     delete flagSpan;
486     time_dgesl.Stop();
487
488     flagSpan = new span(series, 1, _T("time_validate"));
489     Timer time_validate;
490     validate(a, b, x, NSIZE);
491
492     delete flagSpan;
493     time_validate.Stop();
494
495     r.times.push_back({time_allocate.Duration_NS(), time_genRnd.←
        Duration_NS(),
496         time_gauss.Duration_NS(), time_dgesl.Duration_NS(),
497         time_validate.Duration_NS()});
498     // Free the memory
499     for (size_t i = 0; i < NSIZE; ++i) {
500         _aligned_free(a[i]);
501     }
502     _aligned_free(b);
503     _aligned_free(x);
504     _aligned_free(ipivot);
505     // delete[] b;
506     // delete[] x;
507     // delete[] ipivot;
508 }
509 r.CalcAvg();
510 r.PrintToCSV(r.name);
511 time_total.Stop();
512 cout << "Total Time: " << Timer::format(time_total.Duration_NS←
    ());
513 return 0;
514 }
515 }

```